

Udiddit, a social news aggregator

Introduction

Udiddit, a social news aggregation, web content rating, and discussion website, is currently using a risky and unreliable Postgres database schema to store the forum posts, discussions, and votes made by their users about different topics.

The schema allows posts to be created by registered users on certain topics, and can include a URL or a text content. It also allows registered users to cast an upvote (like) or downvote (dislike) for any forum post that has been created. In addition to this, the schema also allows registered users to add comments on posts.

Here is the DDL used to create the schema:

```
CREATE TABLE bad_posts (  
    id SERIAL PRIMARY KEY,  
    topic VARCHAR(50),  
    username VARCHAR(50),  
    title VARCHAR(150),  
    url VARCHAR(4000) DEFAULT NULL,  
    text_content TEXT DEFAULT NULL,  
    upvotes TEXT,  
    downvotes TEXT  
);  
  
CREATE TABLE bad_comments (  
    id SERIAL PRIMARY KEY,  
    username VARCHAR(50),  
    post_id BIGINT,  
    text_content TEXT  
);
```

Part I: Investigate the existing schema

As a first step, investigate this schema and some of the sample data in the project's SQL workspace. Then, in your own words, outline three (3) specific things that could be improved about this schema. Don't hesitate to outline more if you want to stand out!

1. The bad_posts and bad_comments tables are denormalized, leading to redundant and duplicate data.
2. The username field should be unique, non-empty, and each registered user must be linked to a user_id.
3. Topic names created by users must be unique.
4. Upvotes and Downvotes should be stored as INT instead of TEXT since they represent numerical values.
5. In the bad_comments table, the username should be replaced with user_id as a foreign key, referencing the id column in the user table.
6. The post_id should act as a foreign key, linking to the id column in the posts table.
7. The post_id should be of type INT instead of BIGINT.
8. The absence of constraints and indexes in both tables negatively impacts query performance.

Part II: Create the DDL for your new schema

Having done this initial investigation and assessment, your next goal is to dive deep into the heart of the problem and create a new schema for Udiddit. Your new schema should at least reflect fixes to the shortcomings you pointed to in the previous exercise. To help you create the new schema, a few guidelines are provided to you:

1. Guideline #1: here is a list of features and specifications that Udiddit needs in order to support its website and administrative interface:
 - a. Allow new users to register:
 - i. Each username has to be unique
 - ii. Usernames can be composed of at most 25 characters
 - iii. Usernames can't be empty
 - iv. We won't worry about user passwords for this project
 - b. Allow registered users to create new topics:
 - i. Topic names have to be unique.
 - ii. The topic's name is at most 30 characters
 - iii. The topic's name can't be empty
 - iv. Topics can have an optional description of at most 500 characters.
 - c. Allow registered users to create new posts on existing topics:
 - i. Posts have a required title of at most 100 characters
 - ii. The title of a post can't be empty.
 - iii. Posts should contain either a URL or a text content, **but not both**.

- iv. If a topic gets deleted, all the posts associated with it should be automatically deleted too.
 - v. If the user who created the post gets deleted, then the post will remain, but it will become dissociated from that user.
 - d. Allow registered users to comment on existing posts:
 - i. A comment's text content can't be empty.
 - ii. Contrary to the current linear comments, the new structure should allow comment threads at arbitrary levels.
 - iii. If a post gets deleted, all comments associated with it should be automatically deleted too.
 - iv. If the user who created the comment gets deleted, then the comment will remain, but it will become dissociated from that user.
 - v. If a comment gets deleted, then all its descendants in the thread structure should be automatically deleted too.
 - e. Make sure that a given user can only vote once on a given post:
 - i. Hint: you can store the (up/down) value of the vote as the values 1 and -1 respectively.
 - ii. If the user who cast a vote gets deleted, then all their votes will remain, but will become dissociated from the user.
 - iii. If a post gets deleted, then all the votes for that post should be automatically deleted too.
2. Guideline #2: here is a list of queries that Udiddit needs in order to support its website and administrative interface. Note that you don't need to produce the DQL for those queries: they are only provided to guide the design of your new database schema.
- a. List all users who haven't logged in in the last year.
 - b. List all users who haven't created any post.
 - c. Find a user by their username.
 - d. List all topics that don't have any posts.
 - e. Find a topic by its name.
 - f. List the latest 20 posts for a given topic.
 - g. List the latest 20 posts made by a given user.
 - h. Find all posts that link to a specific URL, for moderation purposes.
 - i. List all the top-level comments (those that don't have a parent comment) for a given post.
 - j. List all the direct children of a parent comment.
 - k. List the latest 20 comments made by a given user.
 - l. Compute the score of a post, defined as the difference between the number of upvotes and the number of downvotes

3. Guideline #3: you'll need to use normalization, various constraints, as well as indexes in your new database schema. You should use named constraints and indexes to make your schema cleaner.
4. Guideline #4: your new database schema will be composed of five (5) tables that should have an auto-incrementing id as their primary key.

Once you've taken the time to think about your new schema, write the DDL for it in the space provided here:

```
CREATE TABLE "users" (  
  user_id SERIAL Primary Key,  
  username VARCHAR(25) CONSTRAINT "required_unique_name" UNIQUE NOT NULL,  
  created_user_at TIMESTAMP WITH TIME ZONE  
);  
CREATE INDEX "LOG_IN_INDEX" ON "users" ("created_user_at");  
  
CREATE TABLE "topics"(  
  topic_id SERIAL PRIMARY KEY,  
  user_id INTEGER REFERENCES "users",  
  topic_name VARCHAR(30) CONSTRAINT "REQUIRED_UNIQUE_TOPICNAME" UNIQUE NOT  
  NULL,  
  topic_desc VARCHAR(500)  
);  
  
CREATE TABLE "posts"(  
  post_id SERIAL PRIMARY KEY,  
  topic_id INTEGER REFERENCES "topics" ON DELETE CASCADE,  
  user_id INTEGER REFERENCES "users" ON DELETE SET NULL,  
  created_post_at TIMESTAMP WITH TIME ZONE,  
  post_title VARCHAR(100) NOT NULL,  
  url_content VARCHAR(5000),  
  text_content TEXT,  
  CONSTRAINT "URL_OR_TEXT" CHECK(("url_content" IS NOT NULL AND  
  "text_content" IS NULL) OR ("url_content" IS NULL AND "text_content" IS NOT  
  NULL))  
);  
  
CREATE INDEX "LATEST_TOPICS_INDEX" ON "posts" ("topic_id",  
  "created_post_at");  
CREATE INDEX "TOPICS_PER_USER" ON "posts"("topic_id", "user_id");  
CREATE INDEX "URL_CONTENT" ON "posts"(url_content VARCHAR_PATTERN_OPS);
```

```

CREATE TABLE "comments"(
comment_id SERIAL PRIMARY KEY,
user_id INTEGER REFERENCES "users" ON DELETE SET NULL,
post_id INTEGER REFERENCES "posts" ON DELETE CASCADE,
comment_content TEXT NOT NULL,
created_comment_at TIMESTAMP WITH TIME ZONE,
level INTEGER REFERENCES "comments" ON DELETE CASCADE,
parent_comment_id INTEGER REFERENCES "comments" ON DELETE CASCADE );

CREATE INDEX "LEVEL_INDEX" ON "comments"("level");
CREATE INDEX "LATEST_COMMENTS_PER_USER" ON "comments"("user_id",
"created_comment_at");
CREATE INDEX "COMMENTS_POSTS_INDEX" ON "comments"("post_id");

CREATE TABLE "votes"(
vote_id SERIAL PRIMARY KEY,
user_id INTEGER REFERENCES "users" ON DELETE SET NULL,
post_id INTEGER REFERENCES "posts" ON DELETE CASCADE,
vote INTEGER CONSTRAINT "UPVOTE_DOWNVOTE" CHECK (vote=1 OR vote=-1),
CONSTRAINT "one_vote_per_user" UNIQUE ("user_id", "post_id")
);

CREATE INDEX "VOTE_INDEX" on "votes"("vote");

```

Part III: Migrate the provided data

Now that your new schema is created, it's time to migrate the data from the provided schema in the project's SQL Workspace to your own schema. This will allow you to review some DML and DQL concepts, as you'll be using INSERT...SELECT queries to do so. Here are a few guidelines to help you in this process:

1. Topic descriptions can all be empty
2. Since the bad_comments table doesn't have the threading feature, you can migrate all comments as top-level comments, i.e. without a parent
3. You can use the Postgres string function **regexp_split_to_table** to unwind the comma-separated votes values into separate rows
4. Don't forget that some users only vote or comment, and haven't created any posts. You'll have to create those users too.
5. The order of your migrations matter! For example, since posts depend on users and topics, you'll have to migrate the latter first.
6. Tip: You can start by running only SELECTs to fine-tune your queries, and use a LIMIT to avoid large data sets. Once you know you have the correct query, you can then run your full INSERT...SELECT query.

7. **NOTE:** The data in your SQL Workspace contains thousands of posts and comments. The DML queries may take at least 10-15 seconds to run.

Write the DML to migrate the current data in bad_posts and bad_comments to your new database schema:

```
INSERT INTO "users" ("username")
SELECT DISTINCT username from bad_posts;

INSERT INTO "users" ("username")
select distinct bc.username from bad_comments bc
left join users u
on bc.username=u.username
where u.username IS NULL;

INSERT INTO "users" ("username")
WITH table1 AS (
select regexp_split_to_table("upvotes",'(',')' as upvote_users
from bad_posts)
select distinct upvote_users
from table1 ti
left join users u
on ti.upvote_users=u.username
where u.username IS NULL;

INSERT INTO "topics" ("topic_name")
select distinct topic from bad_posts

INSERT INTO
"posts" ("post_id", "topic_id", "user_id", "post_title", "url_content",
"text_content")

select bp.id, t.topic_id, u.user_id,
LEFT(bp.title,100), bp.url, bp.text_content

from bad_posts bp
join topics t on bp.topic=t.topic_id
join users u on bp.username=u.username;
```

```
INSERT INTO "comments" ("user_id", "post_id", "comment_content",  
"level")  
select u.user_id,p.post_id, bc.text_content, row_number()  
over(partition by p.post_id)  
from bad comments bc  
Join posts pon bc.post_id=p.post_id  
Join users u on be.username=u.username;
```

```
INSERT INTO "votes" ("user_id","post_id", "vote")  
WITH table1 as (bad_posts)  
select id, regexp_split_to_table("upvotes",',') as upvote from  
bad_posts)
```

```
select u.user_id, t1.id, 1 as upvote  
from table1 t1  
join users u  
on t1.upvote=u.username;
```

```
INSERT INTO "votes" ("user_id","post_id","vote")  
WITH table1 as (  
select id, regexp_split_to_table("downvotes",',') as downvote from  
bad_posts)
```

```
select u.user_id, t1.id, -1 as downvote  
from table1 t1  
join users u  
on t1.downvote=u.username;
```